



mongoDB

NoSQL

NoSQL Data Models

- Key-value
- Graph database
- Document-oriented
- Column family

Cloud Databases



Motivations

► Problems with SQL

- Rigid schema
- Not easily scalable (designed for 90's technology or worse)
- Requires unintuitive joins

► Perks of mongoDB

- Easy interface with common languages (Java, Javascript, PHP, etc.)
- DB tech should run anywhere (VM's, cloud, etc.)
- Keeps essential features of RDBMS's while learning from key-value noSQL systems

Introduction

- A Schema-less / Document Oriented Database
 - Data is stored in documents, not tables / relations
- MongoDB is Implemented in C++ for best performance
- **Platforms** : 32/64 bit Windows Linux, Mac OS-X, FreeBSD, Solaris
- **Language drivers** : Ruby/Ruby-on-Rails, Java, C#, JavaScript , C / C++ ,Erlang,Python, Perl...
- Replication & High Availability
- Map/Reduce
- Querying & Fast In-Place Updates

Why use MongoDB

- Simple queries
- Makes sense with most web applications
- Easier and faster integration of data
- Not well suited for heavy and complex transactions systems.
- **Performance / Scalability / Availability**
 - No Joins + No multi-row transactions
 - Fast Reads / Writes
 - Async writes
 - you don't wait for inserts to complete
 - Secondary Indexes
 - Index on embedded document fields for superfast ad-hoc queries

Document store : analogy wrt RDBM

RDBMS		MongoDB
Database	⇒	Database
Table, View	⇒	Collection
Row	⇒	Document (JSON, BSON)
Column	⇒	Field
Index	⇒	Index
Join	⇒	Embedded Document
Foreign Key	⇒	Reference
Partition		Shard

MongoDB's Data Model

- A MongoDB instance may have zero or more databases
 - A Database has “Collections”
 - Collections have “Documents”
 - Documents have “Fields”
 - Fields are key = value pairs
 - A Collection does not enforce the structure of its documents*
- *i.e. Schemaless

JSON

- “JavaScript Object Notation”
- Easy for humans to write/read, easy for computers to parse/generate
- Objects can be nested
- Built on
 - name/value pairs
 - Ordered list of values

BSON

- “Binary JSON”
- Binary-encoded serialization of JSON-like docs
- Also allows “referencing”
- Embedded structure reduces need for joins
- Goals
 - Lightweight
 - Traversable
 - Efficient (decoding and encoding)

BSON Example

```
{  
  "_id" :      "37010"  
  "city" :      "ADAMS",  
  "pop" :      2660,  
  "state" :     "TN",  
  "councilman" : {  
    name: "John Smith"  
    address: "13 Scenic Way"  
  }  
}
```

BSON Types

Type	Number
Double	1
String	2
Object	3
Array	4
Binary data	5
Object id	7
Boolean	8
Date	9
Null	10
Regular Expression	11
JavaScript	13
Symbol	14
JavaScript (with scope)	15
32-bit integer	16
Timestamp	17
64-bit integer	18
Min key	255
Max key	127

The number can be used with the \$type operator to query by type!

The `_id` Field

- By default, each document contains an `_id` field. This field has a number of special characteristics:
 - Value serves as primary key for collection.
 - Value is unique, immutable, and may be any non-array type.
 - Default data type is `ObjectId`, which is “small, likely unique, fast to generate, and ordered.” Sorting on an `ObjectId` value is roughly equivalent to sorting on creation time.

mongoDB vs. SQL

mongoDB	SQL
Document	Tuple
Collection	Table/View
PK: _id Field	PK: Any Attribute(s)
Uniformity not Required	Uniform Relation Schema
Index	Index
Embedded Structure	Joins
Shard	Partition

Installation and Running MongoDB

1. Download from mongodb.org

2. Unzip

3. Create **data directory**

```
>mkdir c:\data\db
```

4. Run MongoDB (**mongod**):

```
>cd c:\mongodb-1.6.3\bin
```

```
>mongod
```

5. Run Mongo shell (**mongo**):

```
>mongo
```

The Mongo Shell

>mongo

>help()

>show dbs

>use <dbname>

>show collections

>db.collectionName.findOne()

>db.collectionName.find()

>db.help()

>db.collectionName.help()

```
C:\appservers\mongo-1.6.3\bin\mongo.exe
MongoDB shell version: 1.6.3
connecting to: test
>
> show dbs
admin
cfmongodb_tests
default_db
local
mongorocks
test
>
>
> use mongorocks
switched to db mongorocks
>
> show collections
people
system.indexes
>
> db.people.findOne()
{
  "_id" : ObjectId("4cb66dae636ac4fa2045ff31"),
  "COUNTER" : NumberLong(1),
  "LOVESMONGO" : true,
  "NAME" : "Marc",
  "BIKE" : "Felt",
  "LOVESSQL" : true,
  "KIDS" : [
    {
      "NAME" : "Alexis",
      "AGE" : NumberLong(7),
      "DESCRIPTION" : "crazy",
      "HAIR" : "blonde"
    },
    {
      "NAME" : "Sidney",
      "AGE" : NumberLong(2),
      "DESCRIPTION" : "ornery",
      "HAIR" : "dirty blonde"
    }
  ],
  "WIFE" : "Heather",
  "TS" : "Wed Oct 13 2010 22:40:46 GMT-0400 (Eastern Daylight Time)"
}
```

CRUD

- Create
 - `db.collection.insert(<document>)`
 - `db.collection.save(<document>)`
 - `db.collection.update(<query>, <update>, { upsert: true })`
- Read
 - `db.collection.find(<query>, <projection>)`
 - `db.collection.findOne(<query>, <projection>)`
- Update
 - `db.collection.update(<query>, <update>, <options>)`
- Delete
 - `db.collection.remove(<query>, <justOne>)`

CRUD example

```
> db.user.insert({  
  first: "John",  
  last : "Doe",  
  age: 39  
})
```

```
> db.user.find ()  
{  
  "_id" : ObjectId("51..."),  
  "first" : "John",  
  "last" : "Doe",  
  "age" : 39  
}
```

```
> db.user.update(  
  {"_id" : ObjectId("51...")},  
  {  
    $set: {  
      age: 40,  
      salary: 7000}  
  }  
)
```

```
> db.user.remove(  
  "first": /^J/  
)
```