

Neo4j


Causal Clustering

Neo4j Causal Clustering

- A cluster is composed of three or more Neo4j instances that communicate with one another to provide fault-tolerance and high-availability
- Uses a **consensus protocol (RAFT)** to coordinate the cluster
- each database has a perfect, complete copy of the entire database (no partitioning)
- Each machine in the cluster has a “**role**” –
 - **Leader** or
 - **Follower**

Cluster Architecture

Neo4j Causal Cluster



Cluster Roles

- The **leader** is responsible for coordinating the cluster and accepting all writes
- **Followers** help scale the read workload ability of the cluster and provide for high-availability of data
- If any one follower fails, show continue
- can have any number caches in the form of read replicas

Topology changes

- In the lifecycle of a cluster, cluster roles are temporary.
- Suppose you have machines A, B, and C.
- If A fails, then the remaining nodes (B and C) will elect a new leader amongst themselves.
- When A restarts, later on, it will rejoin the cluster, but probably as a follower.
- Roles can change through the lifecycle of the cluster
- Role changes are not a cause for concern

Driver API consists of 4 key parts

Driver

Top-level object for all Neo4j interaction

Session

Logical context for sequence of transactions

Transaction

Unit of work

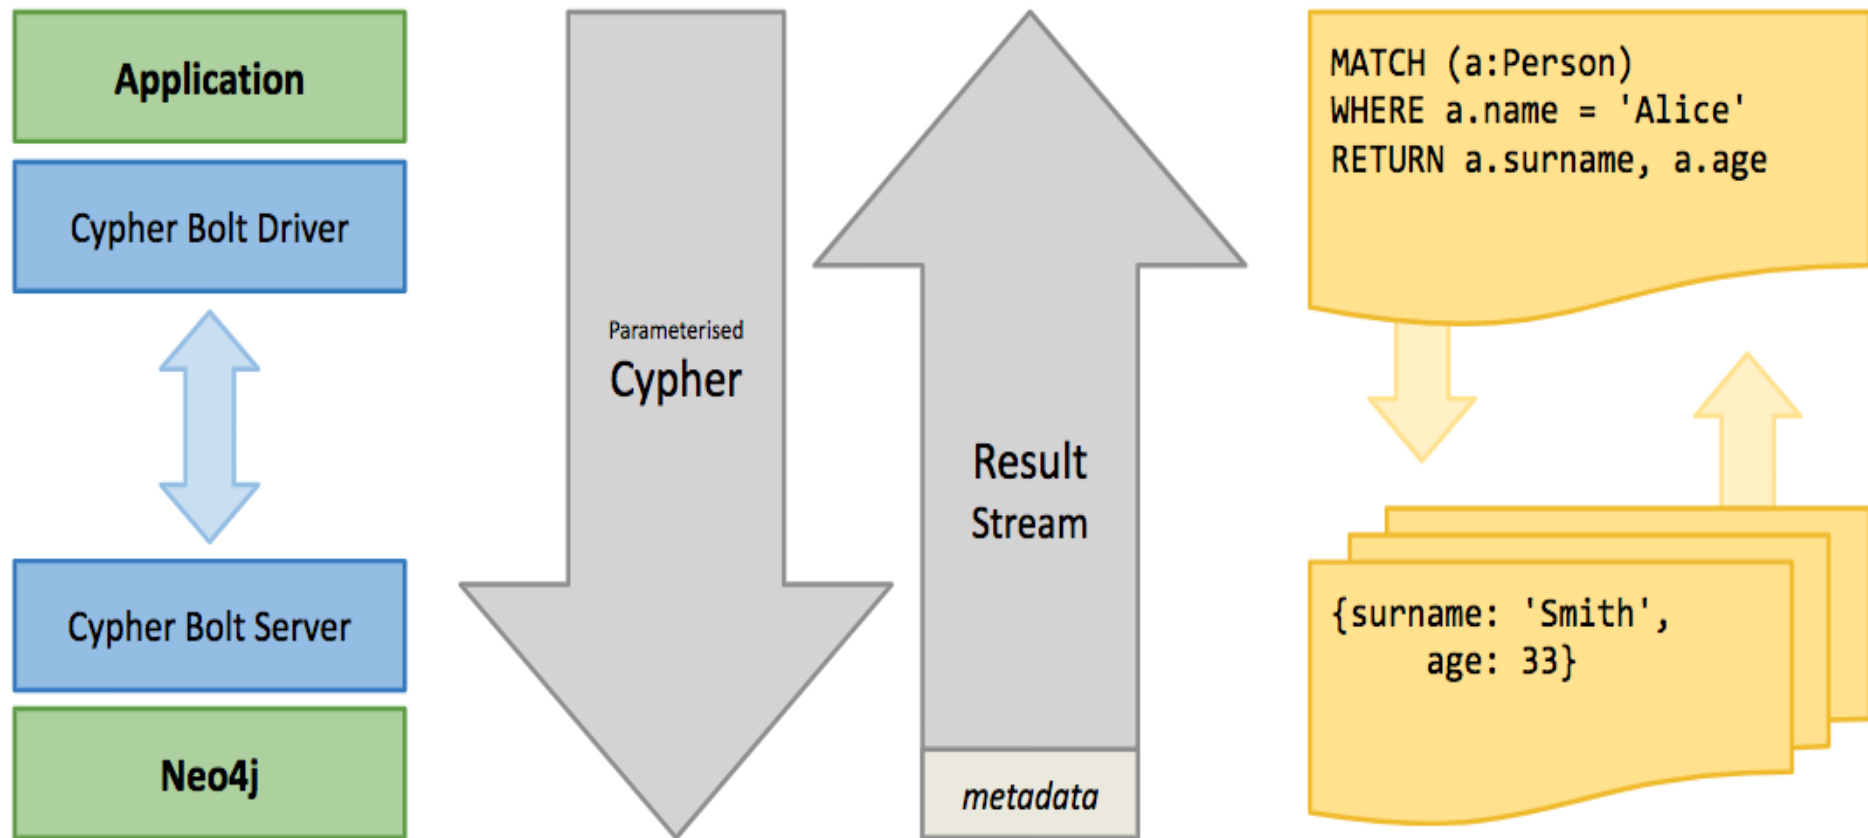
Statement Result

Stream of records plus metadata

Routing Drivers

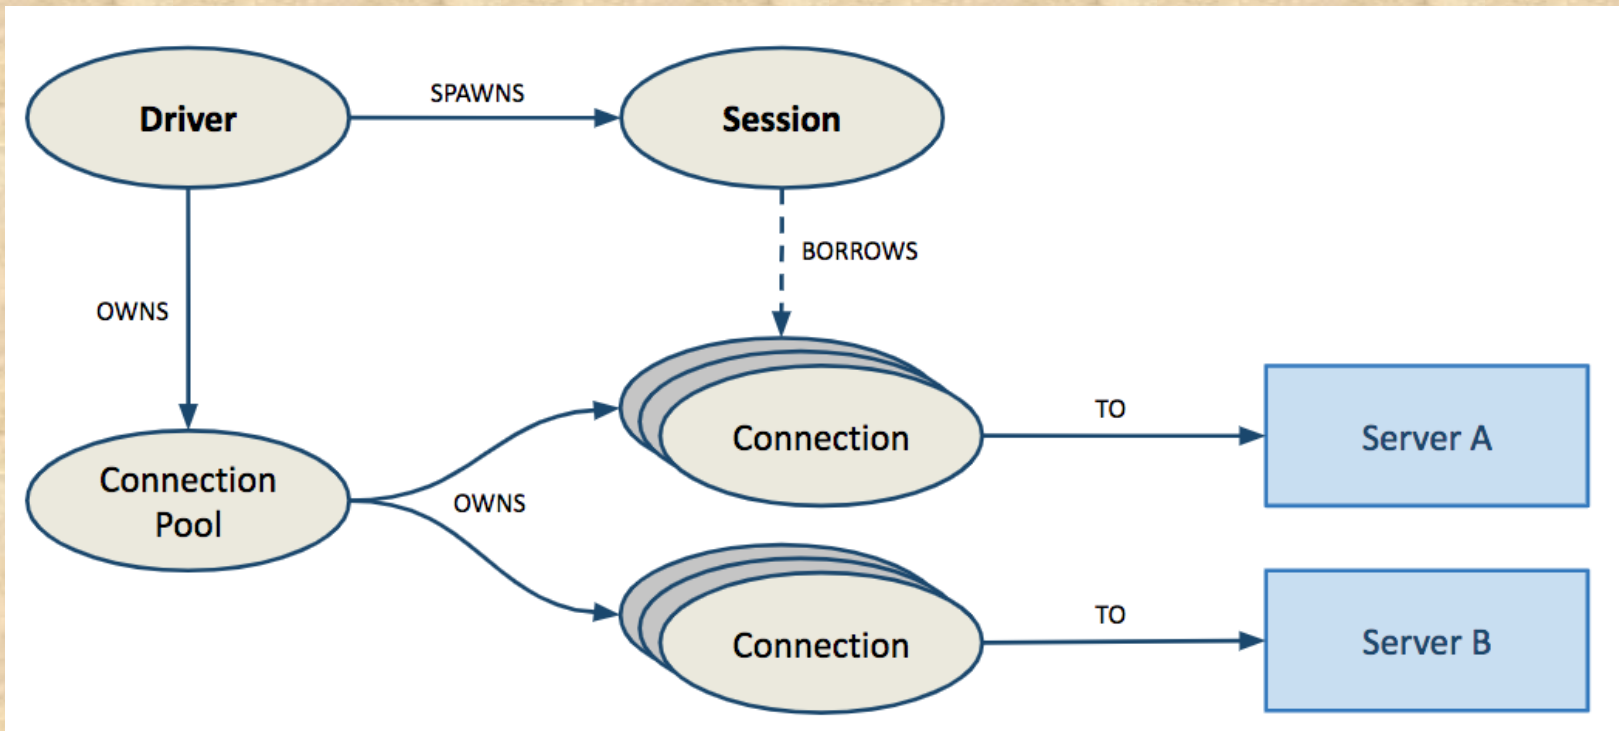
- In one of the supported Neo4j drivers (Java, Javascript, Python, .Net and Go)
- there is an option to use the **bolt+routing** protocol
- E.g. URI of connection string
bolt+routing://neo4j.myhost.com
- Routing driver decides how / where to execute the transactions/queries

How queries get run on Neo4j



Connection Management

- driver manages a **pool of connections** to all machines in the cluster
- user needs to just create the sessions, and run the queries from those sessions !



Installation of Neuo4j Causal Cluster on single machine

Steps

1. Download the enterprise server release from the following link
<http://neo4j.com/download/other-releases/#releases>
2. Copy the downloaded compressed file in 3 separate directories that effectively creates 3 instances
3. Configure *neo4j.conf* file in each instance directory as follows

Modifications in *Neo4j.conf* for each instance

- dbms.backup.enabled=true
- dbms.backup.address=127.0.0.1:6362
- dbms.connector.bolt.address=127.0.0.1:7687
- dbms.connector.http.address=127.0.0.1:7474
- dbms.connector.https.address=127.0.0.1:7473
- dbms.mode=HA
- ha.server_id=1
- ha.initial_hosts=127.0.0.1:5001,127.0.0.1:5002,127.0.0.1:5003
- ha.host.coordination=127.0.0.1:5001
- ha.host.data=127.0.0.1:6001

Steps ...

4. Start up each Neo4j instance with the following commands

/Home directory instance1/bin/Neo4j

/Home directory instance2/bin/Neo4j

/Home directory instance3/bin/Neo4j

5. Startup Neo4J web based admin console in your Browser

http://127.0.0.1:7474 (instance one, HTTP)

http://127.0.0.1:7475 (instance two, HTTP)

http://127.0.0.1:7476 (instance three, HTTP)

Steps ...

6. View the status of your cluster on the monitoring and metrics page of the Neo4J administration console

Store Sizes	ID Allocation	Page Cache	Transactions	High Availability
Array Store 0.00 MB	Node ID 0	Evicts 11	Last Tx Id 1	InstanceId 1
Logical Log 250 B	Property ID 0	Evictions 0	Queued 0	Role master
Node Store 0 B	Relationship ID 0	File Mappings 0	Peak 0	Alive true
Property Store 0 B	Relationship Type ID 0	Bytes Read 121 MB	Opened 11	Available true
Relationship Store 0 B		Flushes 1	Committed 0	Last Committed Tx Id 1
String Store 0.00 MB		Eviction Exceptions 0		Last Update Time 10/11
Total Store Size 130.40 MB		File Unmappings 10		
		Bytes Written 0 MB		

Cluster			
id	Alive	Available	Is Master
1	true	true	yes
2	true	true	-
3	true	true	-